

# Local Orbit Consistency for Symmetry-Robust Neural Operators

Anonymous Authors

## Abstract

Neural operators can learn accurate PDE solution maps while still violating known symmetries of the underlying equations. We study a model-agnostic remedy: stochastic local orbit consistency, a training objective that compares the prediction on a transformed input with the transformed original prediction. The method uses known Lie-group actions, requires no architectural change, supports labeled training and unlabeled adaptation, and leaves inference identical to the base operator. In a periodic N64 2D Navier–Stokes Galilean setting, augmentation plus orbit consistency reduces symmetry-induced OOD error and equivariance defect relative to supervised and augmentation baselines, with the strongest gains in low-label regimes.

## 1 Introduction

Neural operators are attractive for PDE learning because they approximate maps between function spaces rather than single finite-dimensional predictions. This flexibility also exposes a weakness: a high-accuracy operator need not respect symmetries that the exact solution operator satisfies. When test inputs move along a known symmetry orbit, such as a Galilean boost of a periodic flow, ordinary supervised training can produce both higher prediction error and larger equivariance defect.

We propose to reduce this failure mode with local orbit consistency. Given a known group action on inputs and outputs, the objective samples a small finite transform and penalizes the discrepancy between  $\mathcal{G}_\theta(T_g a)$  and  $T_g \mathcal{G}_\theta(a)$ . The penalty is a training-time regularizer, not a new inference pipeline. This makes it complementary to supervised Lie augmentation and architecture-level equivariance: it can be added to an existing FNO, applied in low-label settings, and evaluated on unlabeled inputs.

Our empirical claim is deliberately narrow. On exact periodic 1D and 2D symmetry tests, the coded group actions satisfy solver-level identities. On the main N64 2D Galilean experiment, augmentation plus orbit consistency with the same FNO2d backbone lowers orbit OOD relative  $L^2$  and equivariance defect compared with compute-matched supervised augmentation. The final artifact path, manifest, tables, and figures are regenerated by `python scripts/reproduce_paper_artifacts.py`.

## 2 Related Work

Neural operators learn mappings between function spaces and are therefore a natural model class for parametric PDE solution operators. DeepONet introduced a branch–trunk architecture for learning nonlinear operators from function samples and output coordinates [13]. Fourier Neural Operators (FNOs) parameterize global integral kernels in Fourier space and provide an efficient backbone for PDE families on regular grids [9]. The general neural-operator framework formalizes these models as discretization-invariant compositions of integral operators and nonlinearities, with applications to Burgers’ equation, Darcy flow, and Navier–Stokes [7]. Physics-Informed Neural Operators (PINOs) add PDE residual constraints to operator learning, combining supervised operator regression with physics-based optimization [12]. Subsequent work extends neural operators to geometric settings: Geo-FNO learns deformations from irregular physical domains to latent uniform grids [10], GINO combines graph and Fourier operator components for large-scale three-dimensional PDEs on varying geometries [11], and NORM uses Laplacian eigenfunctions to define neural operators on Riemannian manifolds [3]. Operator-learning benchmarks such as PDEBench and The Well provide reusable datasets and baselines for time-dependent PDEs and broader physical simulation families [14, 17]. This literature establishes FNO-style models as computationally efficient solution-operator backbones, while also showing that domain geometry, discontinuities, and discretization effects require additional structure [8].

Symmetry has been used as an inductive bias through architectures, losses, and data transformations. Group-equivariant convolutional networks replace ordinary convolution by group convolution to share weights across symmetry orbits [4]; steerable CNNs generalize this principle through representation-constrained kernels [19]; and gauge-equivariant CNNs extend equivariance from global transformations to local frame changes on manifolds [5]. For operator learning, Group-Equivariant FNOs encode translation, rotation, and reflection symmetries in Fourier layers [6]. A parallel line derives symmetries from the governing differential equations. Lie Point Symmetry Data Augmentation constructs solution-preserving transformations from a PDE’s Lie point symmetry group and uses them to reduce sample complexity for neural PDE solvers [2]. Lie-symmetry losses for PINNs enforce conservation of solution families under infinitesimal generators [1]. Generalized Lie symmetries for PINOs further argue that evolutionary representatives can provide stronger training signals than point symmetries alone [18]. These works support the central premise that PDE symmetries should constrain the learned solution operator, not only the training data distribution.

Model-agnostic symmetry methods provide the closest precedent for the proposed direction. Tangent propagation used transformation tangents to impose local invariance constraints on neural networks [16]. Lie Algebra Canonicalization uses infinitesimal generators to construct canonicalized inputs for pretrained models and neural operators under arbitrary Lie groups, including non-compact groups commonly arising in PDEs [15]. PACE-FNO applies a related canonicalization strategy to FNOs under symmetry-induced shifts, estimating a Lie-algebraic input frame, predicting in a canonical frame, and restoring the output frame for translations, rotations, and Galilean boosts [20]. Our method studies a complementary point in this design space: a stochastic local orbit-consistency objective that acts directly on an existing neural operator. The regularizer uses known infinitesimal generators, evaluates finite group perturbations, supports labeled training and unlabeled adaptation, and leaves the inference map identical to the base operator.

### 3 Method

We train a neural operator  $\mathcal{G}_\theta : \mathcal{A} \rightarrow \mathcal{U}$  with ordinary supervised regression and a stochastic local orbit-consistency objective. Each minibatch contains input–solution pairs  $(a, u)$ . The supervised term is

$$\mathcal{L}_{\text{sup}}(\theta) = \mathbb{E}_{(a,u)} \left[ \frac{\|\mathcal{G}_\theta(a) - u\|_{\mathcal{U}}^2}{\|u\|_{\mathcal{U}}^2 + \delta} \right]. \quad (1)$$

For a sampled Lie-algebra generator  $X_i$ , scalar step  $\epsilon$ , and finite transform  $g = \exp(\epsilon X_i)$ , the local orbit-consistency term is

$$\mathcal{L}_{\text{orb}}(\theta) = \mathbb{E}_{a,i,\epsilon} \left[ \frac{\|\mathcal{G}_\theta(T_g^{\mathcal{A}} a) - T_g^{\mathcal{U}} \mathcal{G}_\theta(a)\|_{\mathcal{U}}^2}{\epsilon^2 + \eta} \right]. \quad (2)$$

The full objective is  $\min_\theta \mathcal{L}_{\text{sup}}(\theta) + \lambda \mathcal{L}_{\text{orb}}(\theta)$ . For supervised Lie augmentation, we also evaluate  $\mathcal{L}_{\text{aug}}(\theta) = \mathbb{E}_{(a,u),g} [\|\mathcal{G}_\theta(T_g^{\mathcal{A}} a) - T_g^{\mathcal{U}} u\|_{\mathcal{U}}^2 / (\|T_g^{\mathcal{U}} u\|_{\mathcal{U}}^2 + \delta)]$ , and report both augmentation-only and augmentation-plus-orbit variants.

The method is architecture-agnostic. In our experiments,  $\mathcal{G}_\theta$  is an FNO backbone. Transformations are applied to physical-grid input fields; the prediction on the transformed input is compared with the transformed original prediction; gradients pass through both model evaluations. Integer periodic translations are tensor rolls, non-integer periodic translations are spectral shifts, and finite square symmetries are exact tensor rotations and flips. For 1D Burgers, the Galilean action on the solution map  $u_0 \mapsto u_T$  is  $T_c^{\mathcal{A}} u_0(x) = u_0(x) + c$  and  $T_c^{\mathcal{U}} u_T(x) = u_T(x - cT) + c$ .

The same loss supports unlabeled adaptation. Given a pretrained operator  $\mathcal{G}_{\theta_0}$  and target inputs  $a \sim \mathcal{D}_{\text{target}}$ , we minimize  $\mathcal{L}_{\text{orb}}^{\text{target}}(\theta) + \beta \|\theta - \theta_0\|_2^2$  over a chosen adaptation set  $\Theta_{\text{adapt}}$ , such as all parameters, the last operator block, or the projection head. This adaptation objective requires no PDE solves and no target labels.

### 4 Theory

Let  $\mathcal{A}$  and  $\mathcal{U}$  denote Banach spaces of input and output fields. A PDE family induces a solution operator  $\mathcal{S} : \mathcal{A} \rightarrow \mathcal{U}$ . Let  $G$  be a Lie group acting on the full problem family through transformations  $T_g^{\mathcal{A}}$  and  $T_g^{\mathcal{U}}$ .




Figure 1: Local orbit consistency. For  $g = \exp(\epsilon X)$ , let  $a_g = T_g^A a$ ,  $\hat{u} = G_\theta(a)$ ,  $\bar{u}_g = T_g^U \hat{u}$ , and  $\tilde{u}_g = G_\theta(a_g)$ . LOCO penalizes the red equivariance defect  $\Delta_\theta(a, g) = \tilde{u}_g - \bar{u}_g$ , equivalently  $\|\tilde{u}_g - \bar{u}_g\|_{\mathcal{U}}^2/(\epsilon^2 + \eta)$ .

Closure of the full problem class under  $G$  means that the equation, coefficients, forcing terms, boundary or initial conditions, domain class, and observation map transform consistently. Under this assumption the exact solution operator is equivariant:  $\mathcal{S}(T_g^A a) = T_g^U \mathcal{S}(a)$ . For a learned operator  $\mathcal{G}_\theta$ , define the operator-level equivariance defect  $\Delta_\theta(a, g) = \mathcal{G}_\theta(T_g^A a) - T_g^U \mathcal{G}_\theta(a)$ . Our regularizer samples  $g = \exp(\epsilon X_i)$  and penalizes

$$\mathcal{L}_{\text{orb}}(\theta) = \mathbb{E}_{a, i, \epsilon} \left[ \frac{\|\Delta_\theta(a, \exp(\epsilon X_i))\|_{\mathcal{U}}^2}{\epsilon^2 + \eta} \right]. \quad (3)$$

The objective is label-free because both terms are model predictions transformed by known group actions.

**Infinitesimal interpretation.** Assume  $T_{\exp(\epsilon X)}^A$ ,  $T_{\exp(\epsilon X)}^U$ , and  $\mathcal{G}_\theta$  are differentiable at  $\epsilon = 0$ . Let  $X^A a$  and  $X^U u$  denote the induced infinitesimal actions. Taylor expansion gives

$$\mathcal{G}_\theta(T_{\exp(\epsilon X)}^A a) - T_{\exp(\epsilon X)}^U \mathcal{G}_\theta(a) = \epsilon (D\mathcal{G}_\theta(a)[X^A a] - X^U \mathcal{G}_\theta(a)) + O(\epsilon^2). \quad (4)$$

Thus  $\|\Delta_\theta(a, \exp(\epsilon X))\|_{\mathcal{U}}^2/\epsilon^2$  converges to  $\|D\mathcal{G}_\theta(a)[X^A a] - X^U \mathcal{G}_\theta(a)\|_{\mathcal{U}}^2$  as  $\epsilon \rightarrow 0$ . The finite-orbit objective is therefore a stochastic finite-difference estimator of the squared Lie-algebraic equivariance defect, while avoiding explicit Jacobian-vector products.




Figure 2: Vector schematic of local orbit consistency. A local Lie-group step moves an input  $a$  to  $T_g^A a$ . The training penalty compares the prediction on the transformed input,  $G_\theta(T_g^A a)$ , with the transformed original prediction,  $T_g^U G_\theta(a)$ , reducing the equivariance defect along the orbit while leaving inference unchanged.

**Orbit error decomposition.** If the true solution operator is equivariant and  $T_g^U$  is bounded, then for any  $a \in \mathcal{A}$  and  $g \in G$ ,

$$\begin{aligned}
 \|\mathcal{G}_\theta(T_g^A a) - \mathcal{S}(T_g^A a)\|_{\mathcal{U}} &\leq \|\mathcal{G}_\theta(T_g^A a) - T_g^U \mathcal{G}_\theta(a)\|_{\mathcal{U}} \\
 &\quad + \|T_g^U\| \|\mathcal{G}_\theta(a) - \mathcal{S}(a)\|_{\mathcal{U}}.
 \end{aligned} \tag{5}$$

The first term is the learned equivariance defect and the second term is the in-distribution prediction error transported by the output action. This inequality motivates reducing  $\mathcal{L}_{\text{orb}}$  when test inputs differ from training inputs by known PDE symmetries.

**Approximate symmetries and cost.** For partial or approximate symmetries, replace the output norm by a weighted norm  $\|v\|_w^2 = \|w_{a,g} \odot v\|^2$ , where  $w_{a,g}$  can encode valid overlap after transformations, boundary exclusions, observation masks, or regions where transformed coefficients and forcing terms remain physically meaningful. With one sampled transform per minibatch, training uses one ordinary forward pass, one transformed forward pass, and a known output transformation. Inference is exactly the base map  $\mathcal{G}_\theta(a)$ , so deployment latency and memory are unchanged relative to the underlying neural operator.

## 5 Experiments

**Question.** We test whether local orbit consistency reduces symmetry-induced out-of-distribution error and equivariance defect for an unchanged neural-operator backbone. The main experiment uses a periodic  $64 \times 64$  2D Navier–Stokes vorticity operator with Galilean boosts. The base model is the same FNO2d architecture in every row; orbit consistency changes training only and does not add an inference-time module.

Table 1: N64 Galilean headline at 2% labels over five seeds. All rows use the same FNO2d architecture with 2,668,609 parameters. Lower is better for relative  $L^2$ , orbit OOD relative  $L^2$ , and equivariance defect.

| Method             | Steps/epoch | Relative $L^2$                        | Orbit OOD $L^2$                       | Eq. defect                            | Latency ms/sample |
|--------------------|-------------|---------------------------------------|---------------------------------------|---------------------------------------|-------------------|
| FNO                | 4           | $0.9287 \pm 0.0228$                   | $0.9568 \pm 0.0223$                   | $0.6988 \pm 0.0173$                   | $16.53 \pm 2.13$  |
| FNO + orbit        | 4           | $0.4646 \pm 0.0394$                   | $0.4952 \pm 0.0348$                   | $0.3359 \pm 0.0170$                   | $16.11 \pm 1.66$  |
| FNO + aug.         | 6           | $0.5009 \pm 0.0590$                   | $0.5276 \pm 0.0525$                   | $0.3816 \pm 0.0322$                   | $18.76 \pm 7.34$  |
| FNO + aug. + orbit | 4           | <b><math>0.3676 \pm 0.0360</math></b> | <b><math>0.3950 \pm 0.0334</math></b> | <b><math>0.2677 \pm 0.0174</math></b> | $16.87 \pm 3.19$  |

Table 2: Training compute accounting for the headline 2% N64 Galilean runs. Forward/backward counts and transformed-sample counts are computed from the logged optimizer steps; wall-clock minutes are estimated from immutable RNG checkpoint timestamps for these legacy cached runs, and new runs now write `train_wall_seconds` directly. Relative train time is normalized to the compute-matched supervised augmentation baseline.

| Method             | Opt. steps | Model fwds. | Bwd. passes | Aug. labeled samples | Orbit samples | Train min.     | Rel. train time | Latency ms/sample |
|--------------------|------------|-------------|-------------|----------------------|---------------|----------------|-----------------|-------------------|
| FNO                | 600        | 600         | 600         | 0                    | 0             | $14.6 \pm 0.4$ | $0.71\times$    | $16.53 \pm 2.13$  |
| FNO + orbit        | 600        | 1200        | 600         | 0                    | 9600          | $18.7 \pm 0.6$ | $0.91\times$    | $16.11 \pm 1.66$  |
| FNO + aug.         | 900        | 1800        | 900         | 14400                | 0             | $20.6 \pm 2.5$ | $1.00\times$    | $18.76 \pm 7.34$  |
| FNO + aug. + orbit | 600        | 1800        | 600         | 9600                 | 9600          | $21.0 \pm 2.6$ | $1.02\times$    | $16.87 \pm 3.19$  |

**Protocol.** We compare four methods: supervised FNO, supervised Lie augmentation, supervised orbit consistency, and augmentation plus orbit consistency. The headline low-label setting uses 2% labeled data and five seeds. The compute-matched augmentation baseline uses six update steps per epoch, while augmentation plus orbit consistency uses four update steps per epoch with  $\lambda_{\text{orb}} = 0.10$ . All reported rows use the same test-time model architecture and are regenerated by `python scripts/reproduce_paper_artifacts.py`.

**Main N64 Galilean result.** Table 1 shows that at 2% labels, augmentation plus orbit consistency improves both paper-facing symmetry metrics over compute-matched supervised augmentation. Orbit OOD relative  $L^2$  decreases from  $0.5276 \pm 0.0525$  to  $0.3950 \pm 0.0334$ , a 25.1% reduction. Equivariance defect decreases from  $0.3816 \pm 0.0322$  to  $0.2677 \pm 0.0174$ , a 29.8% reduction. Relative to the plain supervised FNO baseline, the reductions are 58.7% for orbit OOD error and 61.7% for equivariance defect. Parameter counts are identical across the headline rows.

**Compute accounting.** Table 2 reports training compute for the same headline runs. The compute-matched augmentation baseline uses 900 optimizer steps and 1800 estimated model forward passes, while augmentation plus orbit consistency uses 600 optimizer steps and the same 1800 estimated model forward passes because each step evaluates the original, augmented, and orbit-transformed inputs. Inference latency remains within run-to-run noise because the deployed map is still the same FNO2d backbone.

**Label efficiency.** Figure 3 and Table 3 show that the benefit persists across the low-label curve. Against compute-matched augmentation,  $\lambda_{\text{orb}} = 0.10$  reduces orbit OOD error by 13.7%, 25.1%, and 22.3% at 1%, 2%, and 5% labels respectively. The corresponding equivariance-defect reductions are 21.3%, 29.8%, and 33.2%. This supports the label-efficiency framing: the orbit objective is most useful when supervision is scarce, while preserving the same inference map.

**OOD severity.** Figure 4 and Table 7 evaluate Galilean boosts from below the training radius to twice the training radius. The augmentation-plus-orbit model maintains a lower orbit OOD error at every tested boost. The relative OOD-error reductions are 26.4%, 25.8%, 25.3%, 23.6%, and 20.3% at maximum boosts 0.10, 0.20, 0.25, 0.35, and 0.50. Equivariance-defect reductions follow the same trend, from 33.6% at boost 0.10 to 23.6% at boost 0.50.

**Canonicalization calibration.** Table 4 reports a privileged relative-boost calibration, not a deployable analytic canonicalization baseline. In the OOD severity protocol, each test input  $a$  is paired with a sampled

Table 3: Label-efficiency comparison for compute-matched augmentation versus augmentation plus orbit consistency with  $\lambda_{\text{orb}} = 0.10$ .

|  | Labels | Method       | Orbit OOD $L^2$                       | Eq. defect                            | Seeds |
|--|--------|--------------|---------------------------------------|---------------------------------------|-------|
|  | 1%     | Aug.         | $0.7421 \pm 0.0838$                   | $0.5080 \pm 0.0456$                   | 3     |
|  | 1%     | Aug. + orbit | <b><math>0.6406 \pm 0.0508</math></b> | <b><math>0.3999 \pm 0.0197</math></b> | 3     |
|  | 2%     | Aug.         | $0.5276 \pm 0.0525$                   | $0.3816 \pm 0.0322$                   | 5     |
|  | 2%     | Aug. + orbit | <b><math>0.3950 \pm 0.0334</math></b> | <b><math>0.2677 \pm 0.0174</math></b> | 5     |
|  | 5%     | Aug.         | $0.3237 \pm 0.0013$                   | $0.2427 \pm 0.0014$                   | 3     |
|  | 5%     | Aug. + orbit | <b><math>0.2515 \pm 0.0112</math></b> | <b><math>0.1621 \pm 0.0078</math></b> | 3     |

Table 4: Privileged relative-boost calibration at 2% labels. Direct OOD evaluates the model on the boosted input. The privileged calibration uses the sampled relative boost from the evaluation protocol and the paired in-distribution prediction, so it is not an observable-frame analytic baseline. Lower is better.

| Max boost | Aug. direct OOD     | Aug. priv. calib.   | Aug. + orbit direct OOD               | Aug. + orbit priv. calib. |
|-----------|---------------------|---------------------|---------------------------------------|---------------------------|
| 0.10      | $0.5050 \pm 0.0574$ | $0.5009 \pm 0.0590$ | <b><math>0.3715 \pm 0.0349</math></b> | $0.3676 \pm 0.0360$       |
| 0.20      | $0.5176 \pm 0.0543$ | $0.5009 \pm 0.0590$ | <b><math>0.3840 \pm 0.0331</math></b> | $0.3676 \pm 0.0360$       |
| 0.25      | $0.5277 \pm 0.0523$ | $0.5009 \pm 0.0590$ | <b><math>0.3943 \pm 0.0318</math></b> | $0.3676 \pm 0.0360$       |
| 0.35      | $0.5555 \pm 0.0476$ | $0.5009 \pm 0.0590$ | <b><math>0.4242 \pm 0.0285</math></b> | $0.3676 \pm 0.0360$       |
| 0.50      | $0.6158 \pm 0.0384$ | $0.5009 \pm 0.0590$ | <b><math>0.4910 \pm 0.0216</math></b> | $0.3676 \pm 0.0360$       |

relative boost  $\delta$ , producing  $T_\delta a$ . The calibration uses this sampled  $\delta$  and the paired in-distribution prediction  $\mathcal{G}_\theta(a)$ , then reports  $T_\delta \mathcal{G}_\theta(a)$ . This is useful for measuring how much of the direct OOD error comes from motion along the tested orbit, but it assumes access to the paired source input and sampled relative transform. It should therefore not be confused with the observable-frame analytic baseline, which reads the absolute ambient velocity  $c$  from the input channels and canonicalizes a single test input. At maximum boost 0.50, augmentation plus orbit consistency has direct OOD error  $0.4910 \pm 0.0216$ , improving substantially over augmentation  $0.6158 \pm 0.0384$  while remaining above its privileged calibration value  $0.3676 \pm 0.0360$ .

**Defect as predictor.** Figure 5 and Table 5 show that equivariance defect is not merely an auxiliary metric. Across 90 cached N64 Galilean observations, defect correlates strongly with orbit OOD relative  $L^2$  (Pearson  $r = 0.851$ , Spearman  $\rho = 0.838$ ); after residualizing source, method, label fraction, maximum boost, and  $\lambda_{\text{orb}}$ , the partial Pearson correlation remains  $r = 0.779$ . The relationship also holds within the OOD severity curve alone ( $r = 0.761$ ), supporting the mechanism that improved orbit consistency tracks improved OOD behavior.

**Ablations and scope.** Table 6 reports the orbit-weight ablation at 2% labels. We keep  $\lambda_{\text{orb}} = 0.10$  as the broader-experiment setting, while the completed high-weight sweep shows  $\lambda_{\text{orb}} = 0.30$  as the strongest tested value. Semi-supervised joint orbit training is useful as an appendix result because it reduces equivariance defect, but it is not the strongest predictive-error result. We do not include a 3D PDE in the main study: a same-resolution 3D analogue would multiply raw field size by 64 and introduce a new solver and validation burden. The current contribution is deliberately narrow: exact-symmetry validation plus a reproducible 2D Galilean low-label result.

## References

- [1] Tara Akhound-Sadegh, Laurence Perreault-Levasseur, Johannes Brandstetter, Max Welling, and Siamak Ravanbakhsh. Lie point symmetry and physics-informed networks. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL [https://proceedings.neurips.cc/paper\\_files/paper/2023/hash/8493c860bec41705f7743d5764301b94-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/8493c860bec41705f7743d5764301b94-Abstract-Conference.html).

Table 5: Correlation between equivariance defect and orbit OOD relative  $L^2$  across cached N64 Galilean seeds, label fractions, orbit weights, and boost severities. Partial Pearson  $r$  residualizes source, method, label fraction, maximum boost, and  $\lambda_{\text{orb}}$ .

| Scope              | Obs. | Pearson $r$ | Spearman $\rho$ | Partial $r$ |
|--------------------|------|-------------|-----------------|-------------|
| All                | 90   | 0.851       | 0.838           | 0.779       |
| Train/lambda/label | 40   | 0.990       | 0.990           | 0.976       |
| OOD severity       | 50   | 0.761       | 0.754           | 0.516       |

Table 6: Orbit-weight ablation at 2% labels.  $\lambda_{\text{orb}} = 0.10$  is used for the broader experiments;  $\lambda_{\text{orb}} = 0.30$  is the strongest among the tested values.

| $\lambda_{\text{orb}}$ | ID $L^2$                              | Orbit OOD $L^2$                       | Eq. defect                            | Seeds |
|------------------------|---------------------------------------|---------------------------------------|---------------------------------------|-------|
| 0.10                   | $0.3676 \pm 0.0360$                   | $0.3950 \pm 0.0334$                   | $0.2677 \pm 0.0174$                   | 5     |
| 0.20                   | $0.3291 \pm 0.0318$                   | $0.3564 \pm 0.0304$                   | $0.2268 \pm 0.0165$                   | 5     |
| 0.30                   | <b><math>0.3149 \pm 0.0298</math></b> | <b><math>0.3416 \pm 0.0289</math></b> | <b><math>0.2044 \pm 0.0160</math></b> | 5     |

- [2] Johannes Brandstetter, Max Welling, and Daniel E Worrall. Lie point symmetry data augmentation for neural PDE solvers. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato, editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 2241–2256. PMLR, 17–23 Jul 2022. URL <https://proceedings.mlr.press/v162/brandstetter22a.html>.
- [3] Gengxiang Chen, Xu Liu, Qinglu Meng, Lu Chen, Changqing Liu, and Yingguang Li. Learning neural operators on riemannian manifolds. *National Science Open*, 3:20240001, 2024. doi: 10.1360/nso/20240001. URL <https://doi.org/10.1360/nso/20240001>.
- [4] Taco Cohen and Max Welling. Group equivariant convolutional networks. In Maria Florina Balcan and Kilian Q. Weinberger, editors, *Proceedings of the 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pages 2990–2999. PMLR, 2016. URL <https://proceedings.mlr.press/v48/cohen16.html>.
- [5] Taco Cohen, Maurice Weiler, Berkay Kicanaoglu, and Max Welling. Gauge equivariant convolutional networks and the icosahedral CNN. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 1321–1330. PMLR, 2019. URL <https://proceedings.mlr.press/v97/cohen19d.html>.
- [6] Jacob Helwig, Xuan Zhang, Cong Fu, Jerry Kurtin, Stephan Wojtowysch, and Shuiwang Ji. Group equivariant Fourier neural operators for partial differential equations. In Andreas Krause, Emma Brunskill, Kyunghyun Cho, Barbara Engelhardt, Sivan Sabato, and Jonathan Scarlett, editors, *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 12907–12930. PMLR, 23–29 Jul 2023. URL <https://proceedings.mlr.press/v202/helwig23a.html>.
- [7] Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces with applications to pdes. *Journal of Machine Learning Research*, 24(89):1–97, 2023. URL <https://www.jmlr.org/papers/v24/21-1524.html>.
- [8] Samuel Lanthaler, Roberto Molinaro, Patrik Hadorn, and Siddhartha Mishra. Nonlinear reconstruction for operator learning of pdes with discontinuities. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=CrfhZAsJDsZ>.
- [9] Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In

Table 7: Relative reductions of augmentation plus orbit consistency over compute-matched augmentation across Galilean OOD severities at 2% labels.

| Max boost | OOD reduction | Defect reduction |
|-----------|---------------|------------------|
| 0.10      | 26.4%         | 33.6%            |
| 0.20      | 25.8%         | 31.3%            |
| 0.25      | 25.3%         | 30.2%            |
| 0.35      | 23.6%         | 27.8%            |
| 0.50      | 20.3%         | 23.6%            |




Figure 3: N64 Galilean label efficiency. Error bars show standard deviation across seeds and faint points show individual seeds. Augmentation plus orbit consistency improves orbit OOD error and equivariance defect at 1%, 2%, and 5% labels.

- International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=c8P9NQVtmn0>.
- [10] Zongyi Li, Daniel Zhengyu Huang, Burigede Liu, and Anima Anandkumar. Fourier neural operator with learned deformations for pdes on general geometries. *Journal of Machine Learning Research*, 24 (388):1–26, 2023. URL <https://www.jmlr.org/papers/v24/23-0064.html>.
- [11] Zongyi Li, Nikola Borislavov Kovachki, Chris Choy, Boyi Li, Jean Kossaifi, Shourya Otta, Mohammad Amin Nabian, Maximilian Stadler, Christian Hundt, Kamyar Azizzadenesheli, and Anima Anandkumar. Geometry-informed neural operator for large-scale 3d pdes. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL [https://proceedings.neurips.cc/paper\\_files/paper/2023/hash/70518ea42831f02afc3a2828993935ad-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/70518ea42831f02afc3a2828993935ad-Abstract-Conference.html).
- [12] Zongyi Li, Hongkai Zheng, Nikola Borislavov Kovachki, David Jin, Haoxuan Chen, Burigede Liu, Andrew Stuart, Kamyar Azizzadenesheli, and Anima Anandkumar. Physics-informed neural operator for learning partial differential equations. *ACM/IMS Journal of Data Science*, 1(3):1–27, 2024. doi: 10.1145/3648506. URL <https://doi.org/10.1145/3648506>.
- [13] Lu Lu, Pengzhan Jin, Guofei Pang, Zhongqiang Zhang, and George Em Karniadakis. Learning nonlinear operators via DeepONet based on the universal approximation theorem of operators. *Nature Machine Intelligence*, 3(3):218–229, 2021. doi: 10.1038/s42256-021-00302-5. URL <https://doi.org/10.1038/s42256-021-00302-5>.
- [14] Ruben Ohana, Michael McCabe, Lucas Thibaut Meyer, Rudy Morel, Fruzsina Julia Agocs, Miguel Beneitez, Marsha Berger, Blakesley Burkhart, Stuart B. Dalziel, Drummond Buschman Fielding, Daniel Fortunato, Jared A. Goldberg, Keiya Hirashima, Yan-Fei Jiang, Rich Kerswell, Suryanarayana Maddu,




Figure 4: OOD severity at 2% labels. The dashed line marks the training boost radius. Augmentation plus orbit consistency remains better than compute-matched augmentation from below the training radius through twice the training radius.

- Jonah M. Miller, Payel Mukhopadhyay, Stefan S. Nixon, Jeff Shen, et al. The well: A large-scale collection of diverse physics simulations for machine learning. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. URL <https://openreview.net/forum?id=00Sx577BT3>.
- [15] Zakhar Shumaylov, Peter Zaika, James Rowbottom, Ferdia Sherry, Melanie Weber, and Carola-Bibiane Schönlieb. Lie algebra canonicalization: Equivariant neural operators under arbitrary lie groups. In *International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=7PLpiVdnUC>.
- [16] Patrice Simard, Bernard Victorri, Yann LeCun, and John Denker. Tangent prop: A formalism for specifying selected invariances in an adaptive network. In *Advances in Neural Information Processing Systems*, volume 4, 1991. URL <https://proceedings.neurips.cc/paper/1991/hash/65658fde58ab3c2b6e5132a39fae7cb9-Abstract.html>.
- [17] Makoto Takamoto, Timothy Praditia, Raphael Leiteritz, Daniel MacKinlay, Francesco Alesiani, Dirk Pflüger, and Mathias Niepert. PDEBench: An extensive benchmark for scientific machine learning. In *Advances in Neural Information Processing Systems*, volume 35, 2022. URL [https://proceedings.neurips.cc/paper\\_files/paper/2022/hash/0a9747136d411fb83f0cf81820d44afb-Abstract-Datasets\\_and\\_Benchmarks.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/0a9747136d411fb83f0cf81820d44afb-Abstract-Datasets_and_Benchmarks.html).
- [18] Amy Xiang Wang, Zakhar Shumaylov, Peter Zaika, Ferdia Sherry, and Carola-Bibiane Schönlieb. Generalized lie symmetries in physics-informed neural operators, 2025. URL <https://arxiv.org/abs/2502.00373>.
- [19] Maurice Weiler and Gabriele Cesa. General  $E(2)$ -equivariant steerable CNNs. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://arxiv.org/abs/1911.08251>.
- [20] Jiaxiao Xu, Changhong Mou, Yeyu Zhang, and Fengxiang He. Physics-aligned canonical equivariant fourier neural operator under symmetry-induced shifts, 2026. URL <https://arxiv.org/abs/2605.18606>.




Figure 5: Equivariance defect predicts Galilean orbit OOD error across cached N64 seeds, label fractions, orbit weights, and boost severities. Circles are trained-model evaluations and squares are severity evaluations.